

# EFFECTIVE PLATFORM BUILDING WITH KUBERNETES.



## IS K8S NEW LINUX?

Wojciech Barczynski - SMACC.io | Hypatos.ai  
Wrzesień 2018

# WOJCIECH BARCZYŃSKI

- Lead Software Engineer  
& System Engineer
- Interests:  
working software
- Hobby:  
teaching software  
engineering



# BACKGROUND

- ML FinTech ➡ micro-services and k8s
- Before:  
1 z 10 Indonesian mobile e-commerce (Rocket Internet)
- Spent 3.5y with Openstack, 1000+ nodes, 21 data centers
- I do not like INFRA :D

# STORY

- Lyke - [12.2016 - 07.2017]
- SMACC - [10.2017 - present]


# KUBERNETES



# WHY?

- Administracja jest trudna i kosztowna
- Virtualne Maszyny, ansible, salt, etc.
- Za dużo ruchomych części
- Nie kończąca się standaryzacja

# MIKROSERWISY AAA!



# WHY?

- Cloud is not so cheap - \$\$\$



# IMAGINE

- do not need to think about IaaS
- no login on a VM
- less gold plating your CI / CD ...
- DC as a black box


# KUBERNETES

- Container management
- Service and application mindset
- Simple Semantic\*
- Independent from IaaS provider

# KUBERNETES


- Batteries for your 12factory apps
- Service discovery, meta-data support
- Utilize resources to nearly 100%

# KUBERNETES




```
make docker_push; kubectl create -f app-srv-dpl.yaml
```

# SCALE UP! SCALE DOWN!



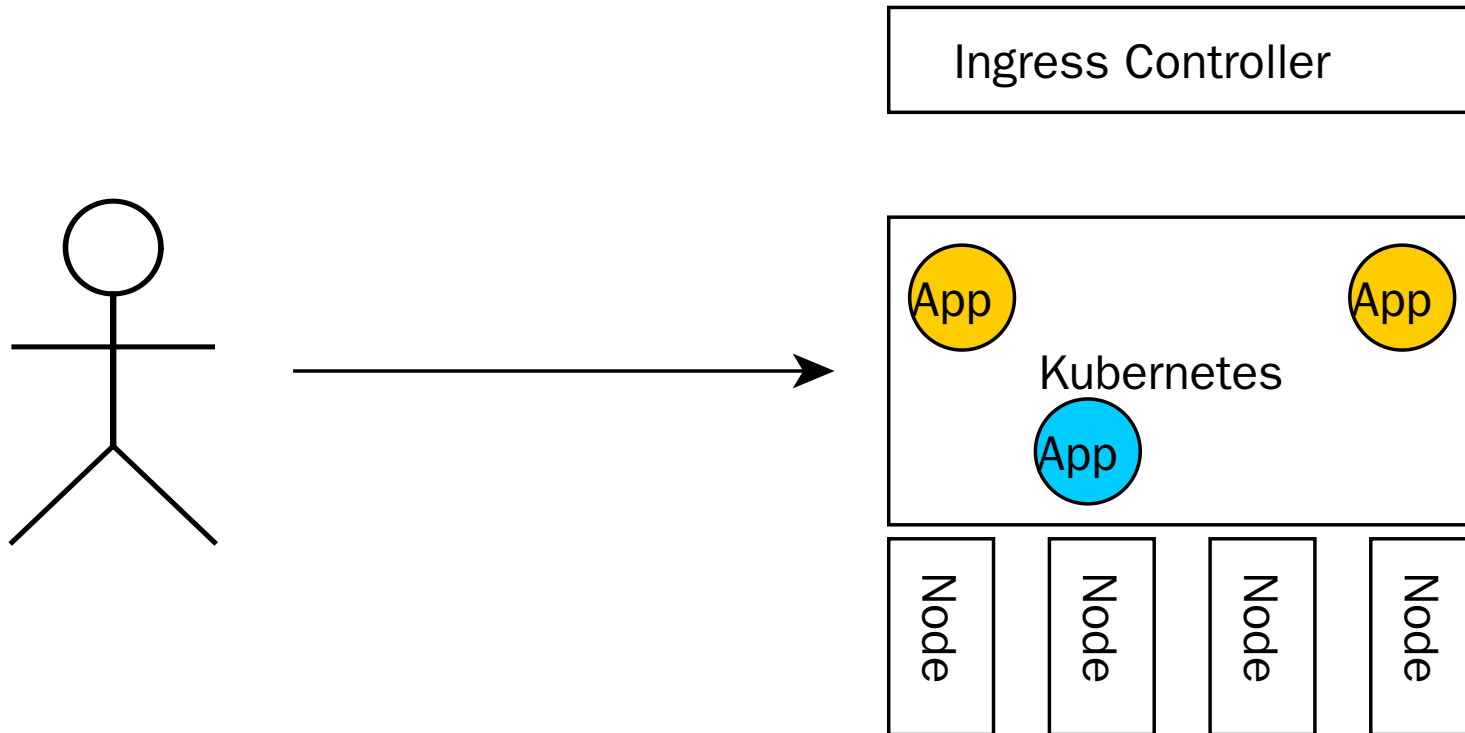
```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

# SCALE UP! SCALE DOWN!



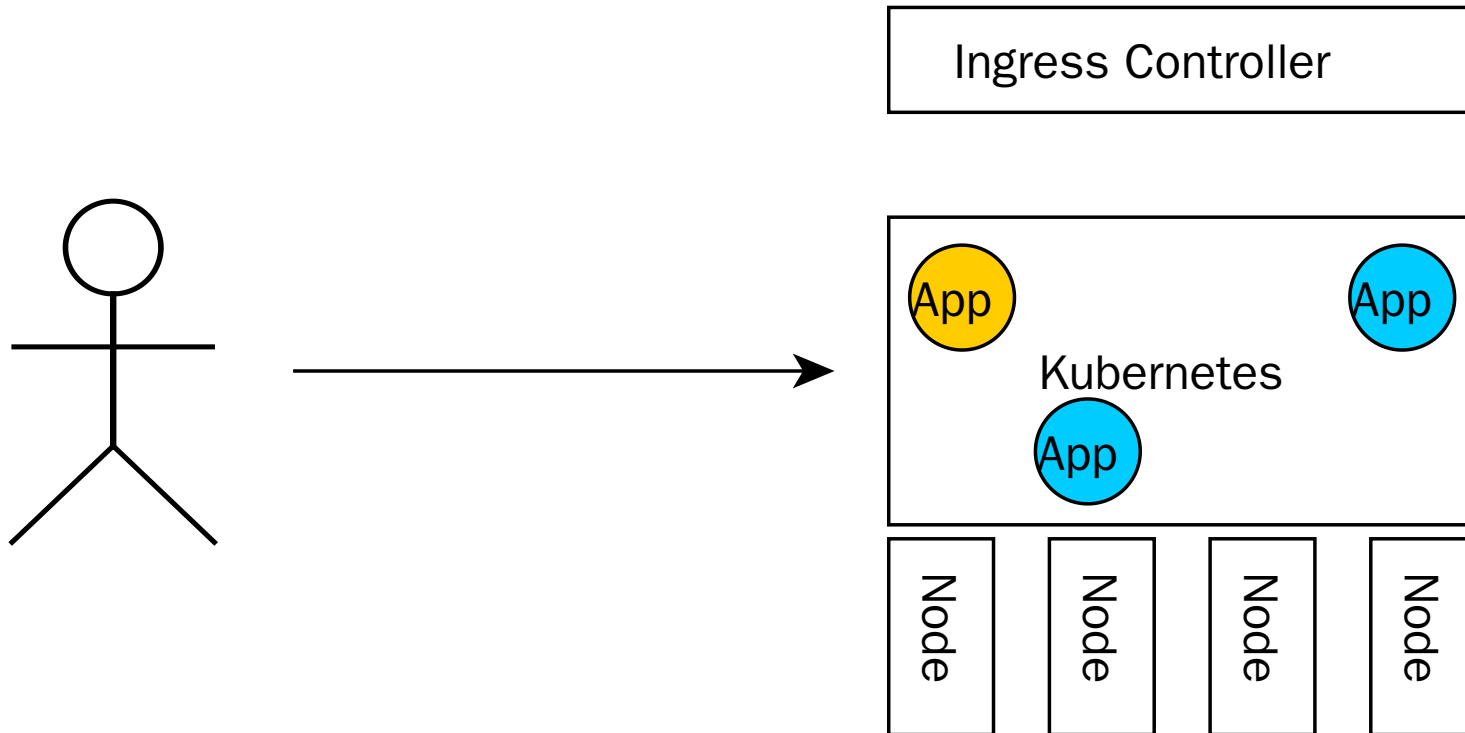
```
kubectl --replicas=1 -f app-srv-dpl.yaml
```

# ROLLING UPDATES!




```
kubectl set image deployment/app app=app:v2.0.0
```

# ROLLING UPDATES!







# ROLLING UPDATES!




# RESISTANCE!




# RESISTANCE!



# RESISTANCE!



# HOW GET USER REQUESTS?




Ingress Controller

# INGRESS

Pattern	Target App Service
api.smacc.io/v1/users	users-v1
api.smacc.io/v2/users	users-v2
smacc.io	web

# LOAD BALANCING



# SERVICE DISCOVERY

- names in DNS:

```
curl http://users/list
```

- labels:

```
name=value
```

- annotations:

```
prometheus.io/scrape: "true"
```



# SERVICE DISCOVERY

- loosely couple components
- auto-wiring with logging and monitoring

# DROP-IN

- traefik / Ingress / Envoy
- prometheus
- audit checks
- ...

# THE BEST PART

All live in git:

- all in Yaml
- integration with monitoring, alarming
- integration with ingress-controller
- ...
- Devs can forget about infrastructure... almost


DevOps Culture Dream!



LYKE

# LYKE

- E-commerce
  - Mobile-only
  - 50k+ users
  - 2M downloads
  - Top 10 Fashion Apps
- w Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

# GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,  
UI Designer,  
Mobile Dev,  
and tester - one  
body



## CHALLENGES

- 50+ VMs in Amazon, 1 VM - 1 App, idle machine
- Puppet, hilarious (manual) deployment process
- Fear
- Forgotten components
- sometimes performance issues

# APPROACH

1. Simplify infrastructure
2. Change the Development practices
3. Change the work organization

see: Conway's law



## **SIMPLIFY**

1. Kubernetes with Google Kubernetes Engine
2. Terraform for all new

## **SIMPLIFY**

1. Prometheus, AlertManager, and Grafana
2. Elasticsearch-Fluentd-Kibana
3. Google Identity-Aware-Proxy to protect all dev dashboards
4. 3rd party SaaS: statuscake and opsgenie

# CONTINUOUS DEPLOYMENT

- branch-based:
  - master
  - staging
  - production
- repo independent

# TRAVIS CI

1. Tests
2. Build docker
3. Deploy to Google Container Registry
4. Deploy to k8s only new docker
5. no config applied

# GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \ - kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```

# Makefile

```
SERVICE_NAME=v-connector
GCP_DOCKER_REGISTRY=eu.gcr.io
test: test_short test_integration

run_local:

docker_build: docker_push

kube_create_config:

kube_apply:

kube_deploy:
```

Copy&Paste from the project to project

# 1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

## 2. GET BACK ALL THE KNOWLEDGE

- Puppet, Chef, ... ➡ Dockerfile
- Check the instances ➡ Dockerfile, README.rst
- Nagios, ... ➡ README.rst, [checks/](#)



### 3. INTRODUCE RUN\_LOCAL

- `make run_local`
- A nice section on how to run in README.rst
- Use: `docker-compose`

The most crucial point.

## 4. GET TO KUBERNETES

- `make kube_create_config`
- `make kube_apply`
- Generate the yaml files if your envs differ

## 5. CONTINUOUS DEPLOYMENT

Travis:

- use the same Makefile as a developer

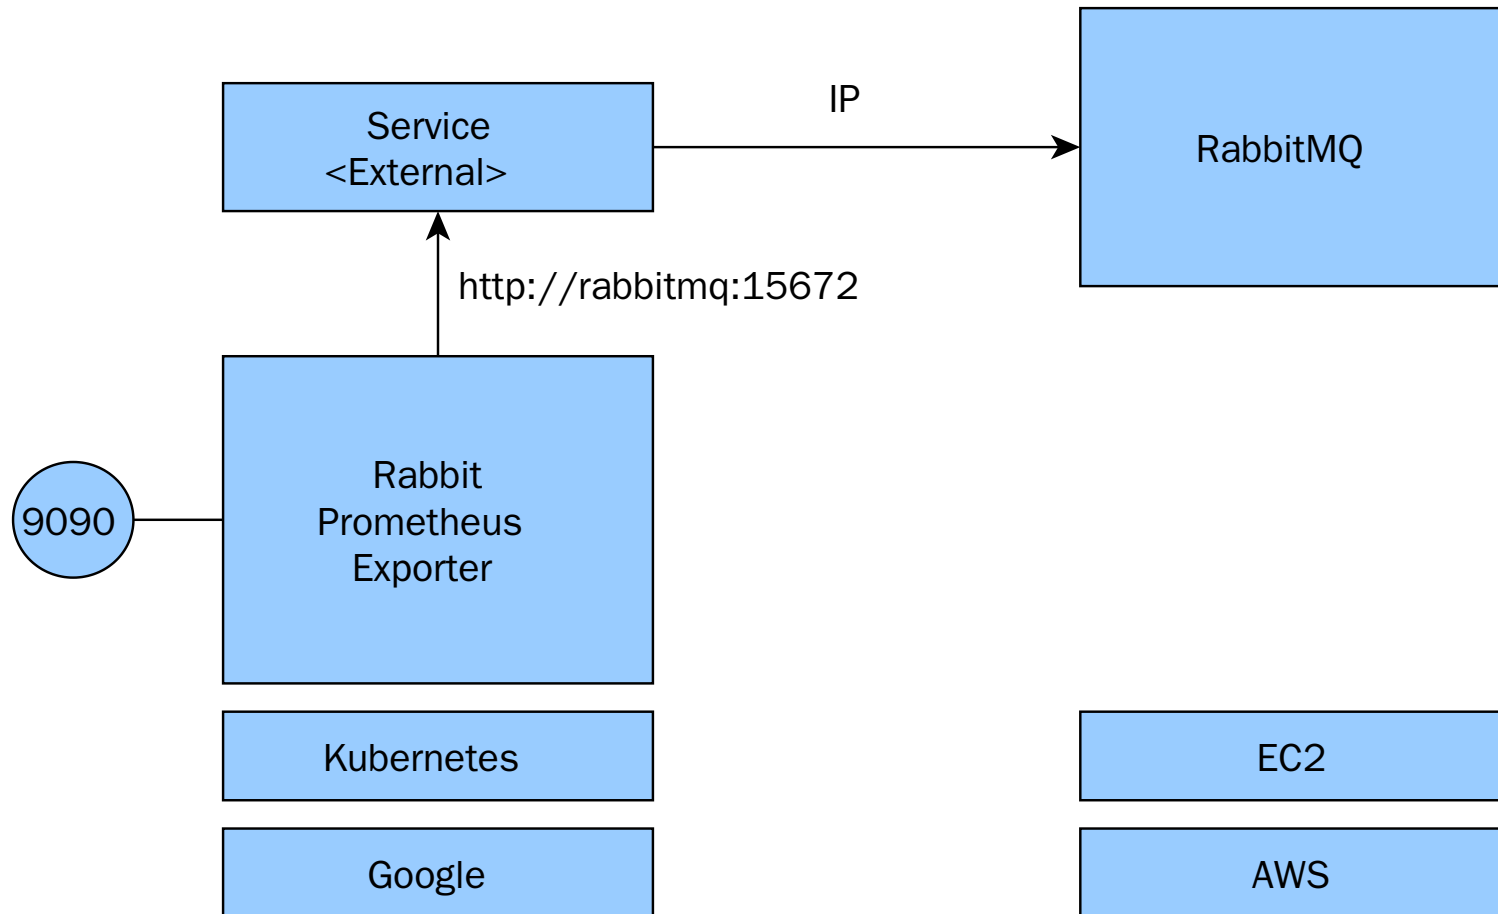
## 6. KEEP IT RUNNING

Bridge the new with old:

- Use external services in Kubernetes
- Optional: Expose k8s in the Legacy [1]


[1] feeding K8S events to HashiCorp consul

# Bridge the new with old



Monitor legacy with new stack

# Architecture During Migration



## 7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test
```

```
TARGET_URL=api.example.com/users make smoke_test
```

## 8. MOVE TO MICRO-SERVICES

To offload the biggest components:

- Keep the lights on
- New functionality delegated to micro-services



## 9. SERVICE SELF-CONSCIOUSNESS

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

## 10. GET PERFORMANCE TESTING

- introduce *wrk* for evaluating performance
- load test the real system

# WHAT WORKED

1. C&P Makefile and k8s between repos
2. Separate deployments a good transition strategy

# WHAT DID NOT WORK

1. Too many PoC, should cut them to 2 weeks max
2. Do it with smaller chunks
3. Alert rules too hard to write
4. Push back to k8s yaml [\*]

With coaching, I thought, it is OK

# DO DIFFERENT

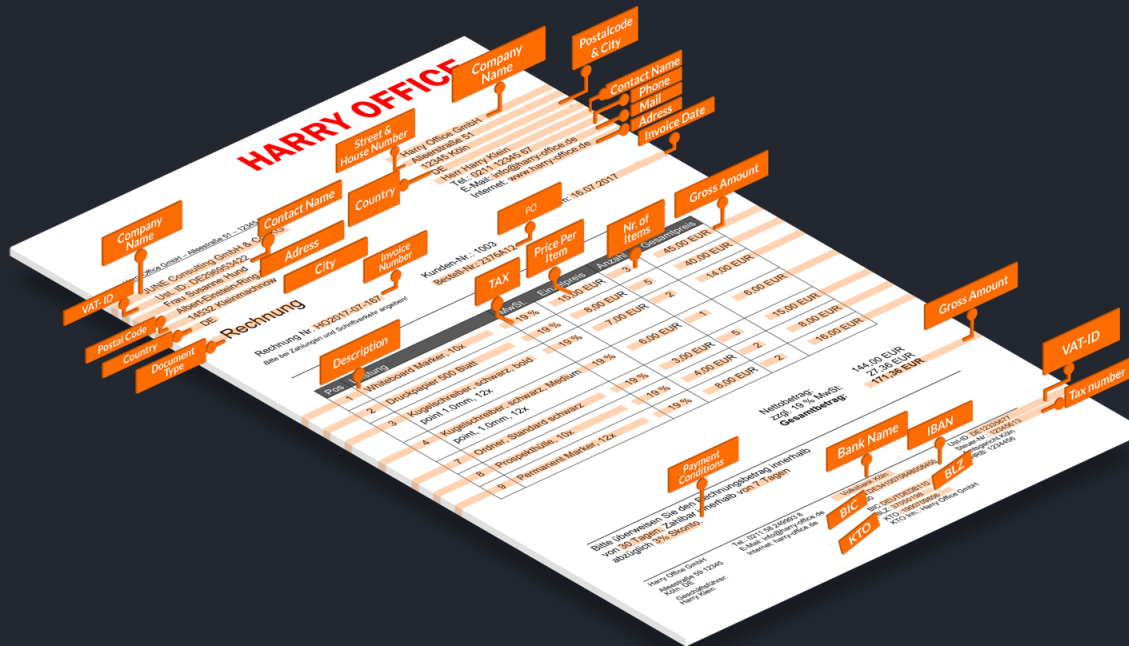
1. Move dev and staging data immediately
2. Let devs know it is a transition stage
3. Teach earlier about resources
4. EFK could wait
5. World-stop for a paid-XXX% weekend for migration

SMACC



Hypatos

## Problem SMACC solves



# Hypatos



SMACC

## Global clients



McKinsey&Company

Hypatos

SMACC



# STORY

- Legacy on AWS, experiments with AWS ECS :/
- Self-hosted K8S on ProfitBricks
- Get to Microsoft ScaleUp, welcome Azure
- Luckily - AKS

# AZURE KUBERNETES SERVICE

- Independent from IaaS
- Our OnPrem = Our OnCloud
- Consolidation of our micro-services
- Plug and play, e.g., monitoring

# SIMPLICITY

- az aks CLI for setting k8s - README.rst
- Terraform for everything else
- 1Password and gopass.pw

TF also sets our AWS

# DIFFERENCE

- Two teams in Berlin and Warsaw
- Me in Warsaw

# NEW EXPERIENCE

- devs really do not like TravisCI ... k8s yamls
- transition from PB to AKS was painful

# SOLUTION

- make everything lighter
- c&p without modifications
- hide the k8s, remove magic
- deploy on *tag*

Similar to the [Kelsey Hightower approach](#)

# Repo .travis.yml

```
language: go
go:
- '1.10'
services:
- docker
install:
- curl -sL https://${GITHUB_TOKEN}@raw.githubusercontent.com
- if [ -f "tools/travis/install.sh" ]; then bash tools/travi
script:
- dep ensure
- make lint
- make test
- if [ -z "${TRAVIS_TAG}" ]; then make snapshot; fi;
deploy:
- provider: script
```

# Makefile

```
| - tools  
|   | - Makefile  
|   | - kube-service.yaml  
|   \- kube-deployment.yaml  
|  
| - Dockerfile  
\- Makefile
```



# CONTINUOUS DEPLOYMENT

- Github
- TravisCI
- [hub.docker.com](https://hub.docker.com)
- AKS

# PROCESS

1. `git tag` and push

# PROCESS

1. Generate deploy, ingress, and svc kubernetes files
2. Commit to smacc-platform.git on **staging** branch
3. Deploy to staging environment

# PROCESS

1. Create PR in smacc-platform.git for **production** branch
2. On merge, deploy to production

# smacc-platform

- 3 independent branches: dev, staging, and master
- Target for other scripts

# KUBERNETES

- Pure, generated, kubernetes config
- 2x kubernetes operators

# WHAT WORKED

- Hiding k8s
- Go for ubuntu-based docker images

# WOULD DO DIFFERENT

- More sensitive to feedback



# NEXT

- Acceptance tests on every deployment
- Scale our ML trainings on the top of k8s
- Deployment tool based on [missy](#)
- Keeping an eye on Istio

# K8S - Linux

- Kubernetes not a silver bullet, but damn close
- Common runtime for onPrem and onCloud
- The biggest asset - the API
- With service discovery - an integration platform
- With kubevirt - might replace your Openstack

DZIĘKUJĘ. PYTANIA?

ps. We are hiring.

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE  
BE WITH YOU.



# BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE  
BE WITH YOU.



# HIRING

- Senior Polyglot Software Engineers
- Experienced System Engineers
- Front-end Engineers
- 1 Data-Driven Product Manager

Apply: [hello-warsaw@smacc.io](mailto:hello-warsaw@smacc.io),  
Questions? [wojciech.barczynski@smacc.io](mailto:wojciech.barczynski@smacc.io), [FB](#) or [LI](#)

We will teach you Go if needed. No k8s or ML, we will take care of that.

**0.1 → 1.0**

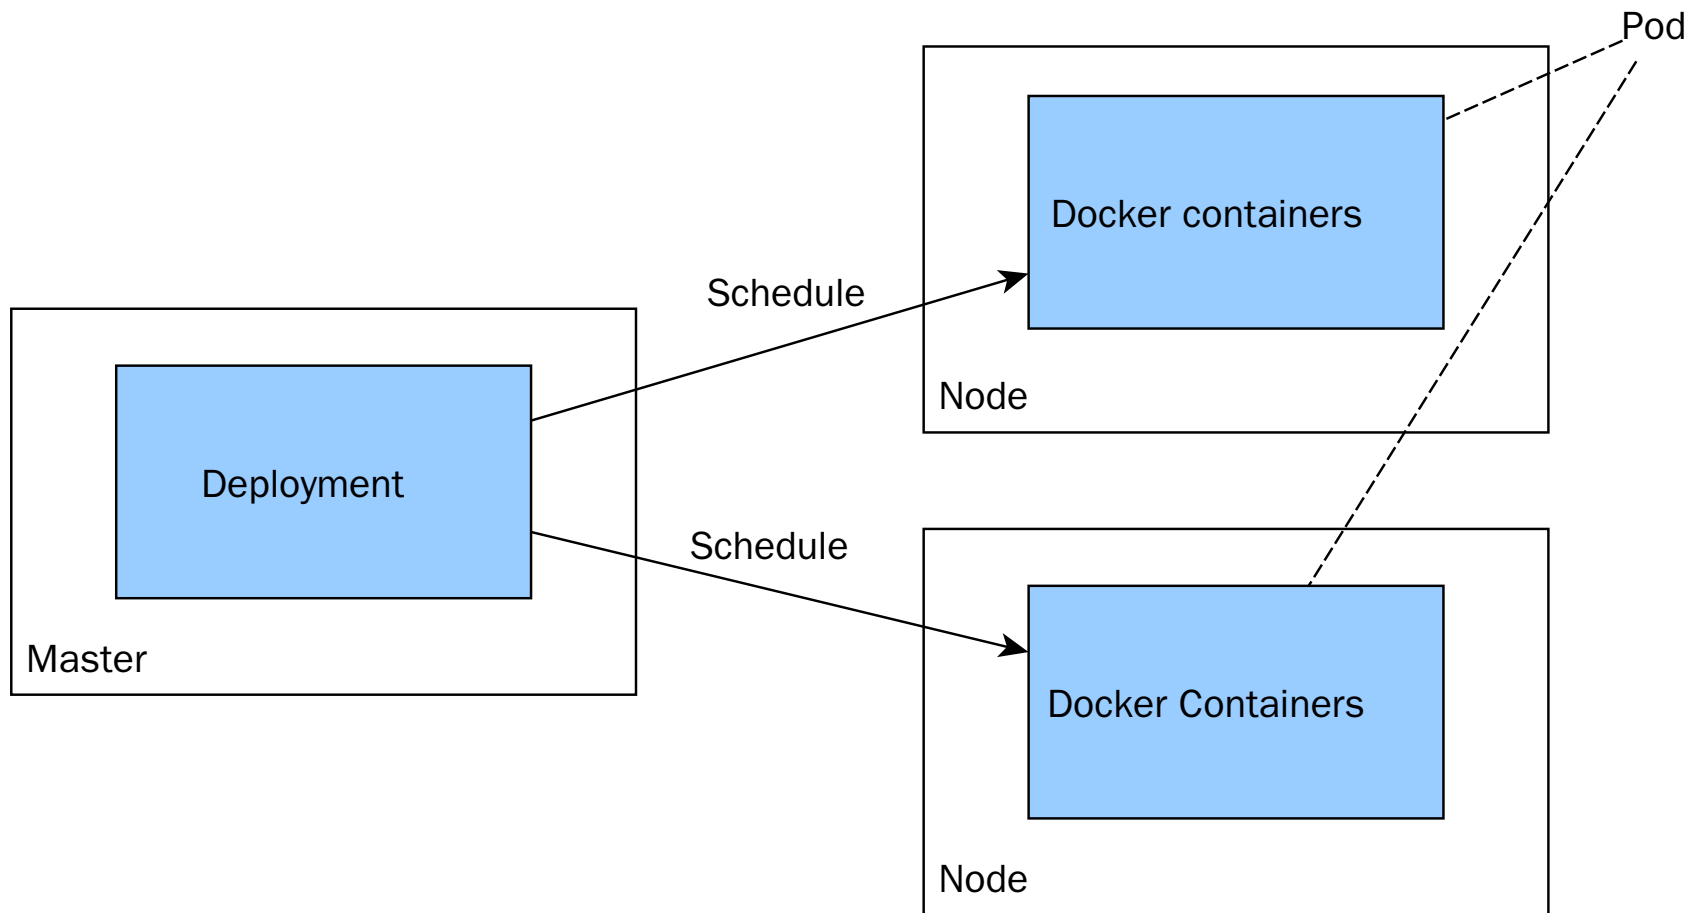
# CHANGE THE WORK ORGANIZATION


- From Scrum
- To Kanban

For the next talk

# KUBERNETES CONCEPTS






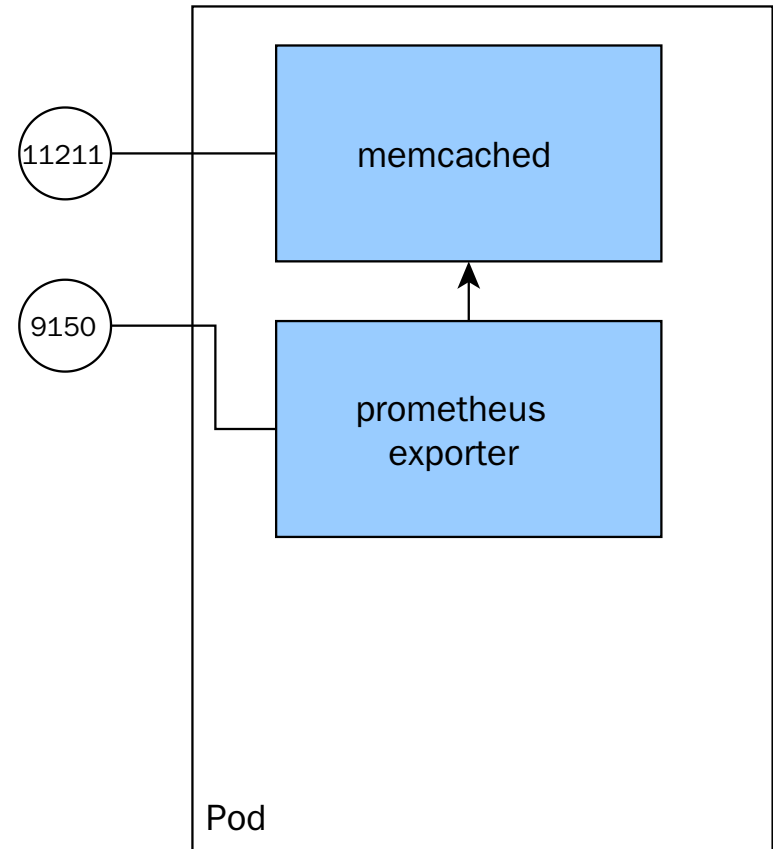



# PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



# SIDE-CARS



## BASIC CONCEPTS

**Name**

**Purpose**

Service

Interface

Entry point  
(Service Name)

---

Deployment

Factory

How many pods,  
which pods


---

Pod

Implementation

1+ docker running

# ROLLING RELEASE WITH DEPLOYMENTS



Also possible